

Beyond Asking: A Pipeline for Personalized Game Generation that Reads Players from Behavior

Yifan Lu¹, Xiaopeng Yuan², Haohan Wang²

¹Imperial College London

²University of Illinois Urbana-Champaign

Abstract

Personalized game generation requires inferring a player’s abilities and behavioral style from how they play. Large language models have made this inference more attainable than ever: an LLM can read a raw gameplay transcript and produce a fluent, plausible profile of the player. Plausible, however, is not verified, and verification is precisely what the field lacks: latent traits are unobservable; questionnaires provide noisy proxies and become circular when self-reports are used to validate behavior-based inference; and behavior itself is ambiguous without context—a player who never collects an item may not want it, or may never have had the chance. We address both problems. First, we construct a synthetic player population whose traits are ground truth by construction: each trait is an explicit bot parameter, accepted only after controlled manipulation produces consistent, trait-specific behavioral change. Unlike prior parameter-recovery work that inverts a known decision model, our benchmark evaluates policy-agnostic inference from behavioral transcripts alone. Second, we introduce an opportunity-aware decision-moment representation that disentangles preference from the chance to express it; ablating it selectively degrades opportunity-dependent traits. On this benchmark, few-shot LLM inference outperforms embedding- and rule-based baselines on most traits, though feature-based supervised regressors remain stronger overall. Finally, we close the loop: inferred profiles drive difficulty adaptation, evaluated against ground-truth references and mismatched-profile controls, and an exploratory human study examines whether these findings transfer to real players.

Introduction

Personalization is a long-standing goal of game content generation (Shaker, Togelius, and Nelson 2016; Yannakakis and Togelius 2018). If a system understands a player’s abilities and behavioral style, it can generate levels with suitable difficulty and compatible gameplay characteristics (Yannakakis and Togelius 2011). The development of large language models (LLMs) makes it feasible to infer latent behavioral tendencies directly from gameplay traces (Gallotta et al. 2024; Peters, Cerf, and Matz 2024). An LLM can read a behavioral transcript and estimate a player’s tendencies through natural-language reasoning, without requiring a separate classifier for each tendency dimension.

However, this approach faces an evaluation problem that precedes any specific modeling method: a player’s latent behavioral tendencies are not directly observable, and their

assessment has long relied on annotation via self-reports or external observers (Yannakakis et al. 2013; Goel, Mahmoudi-Nejad, and Guzdial 2024; Bunian et al. 2017). Whether the inference is performed by rules, supervised learning, or an LLM, it is therefore difficult to determine whether the inferred tendencies are correct. A common practice is to use questionnaires or self-reports as reference labels (John and Srivastava 1999; Zhu, Jin, and Coifman 2025), but self-assessments are noisy and biased (Crowne and Marlowe 1960; Podsakoff et al. 2003; Nisbett and Wilson 1977; Yannakakis and Togelius 2011). Moreover, when the motivation for behavioral inference is precisely to avoid reliance on self-reports, validating such inference against self-reports introduces a potential circularity. The consequence is that fluent player profiles can accumulate faster than the means to falsify any of them: an inference that is wrong in a systematic direction—capping one trait, inflating another—cannot be detected, let alone repaired. Prior work on procedural personas (Holmgård et al. 2014, 2019), inverse reinforcement learning (Ng and Russell 2000), and cognitive modeling (Baker, Saxe, and Tenenbaum 2009; Shergadwala, Teng, and Seif El-Nasr 2021) shares a mature structure for exactly this situation: behavior is generated from known parameters, and an inference procedure is judged by whether it recovers them. Nevertheless, the emerging use of LLMs as behavior-to-tendency inference models still lacks an evaluation environment in which the ground truth is controlled, its validity is independently testable, and the boundaries of inference failure can be identified (Gallotta et al. 2024).

This paper presents an end-to-end personalization pipeline. It also provides a verifiable foundation for the component that has long remained hardest to evaluate: a synthetic diagnostic environment with controlled ground truth. Each behavioral tendency is implemented as an explicit generative parameter in the decision policy of a scripted agent, so the truth is determined by construction rather than by post hoc annotation. A known parameter, however, does not automatically constitute valid ground truth. We therefore introduce a ground-truth admission test: a parameter is accepted only if controlled manipulation produces monotone, dimension-specific behavioral change that exceeds seed-induced variation, measured on held-out behaviors that play no role in defining the parameter; failures are reported as rejected or non-identifiable. This converts recovery evaluation from a

weakly supervised, difficult-to-falsify problem into a controlled and falsifiable one. Its scope is deliberately limited: the environment validates inference methods under controlled conditions rather than substituting for ground truth about real players, and transfer to real players is examined only preliminarily through an exploratory human study.

We instantiate the diagnostic environment in a vertically scrolling shooter. The genre is chosen for measurability: traits are observable only where the situation affords their expression (Tett and Burnett 2003), and this genre supplies such opportunities densely, as discrete decision moments at a fixed rate, each with a choice set the engine already enumerates. At each decision point the record keeps the game state, the set of feasible choices, and the action actually selected, so that behavioral tendency is separated from behavioral opportunity; an LLM reads these records to infer four trait dimensions, and a rule layer maps the inferred profile to generated levels. Nothing in the pipeline is specific to the genre: it transfers to any game whose play can be translated into the same record schema, and what changes across games is the translation layer, not the reading or the validation around it. Our contributions are:

- **A controlled synthetic player population** whose four trait parameters are set by construction, providing the reference ground truth that LLM-based player modeling has so far lacked.
- **An opportunity-aware record** that keeps declined game action choices on file beside taken ones; ablating the choice sets degrades exactly the traits that depend on opportunity.
- **An LLM player model** that reads the opportunity-aware record and infers a four-dimensional trait profile from a fixed example library collected once, with no gradient updates and no access to the generating policy.
- **A closed-loop personalization experiment:** the LLM-inferred profile conditions a generator that assembles a level matched to that player’s traits, and in a pilot with 12 participants, levels generated from the behavioral profile score higher than levels generated from a questionnaire.

Related Work

Inferring latent traits, and what serves as truth. LLMs have been used to infer personality and other latent variables from text or behavioral records (Zhu et al. 2025; Zhu, Jin, and Coifman 2025), but the ground truth in these studies comes almost entirely from self-report scales, the very criterion this paper tries to avoid; psychometric surveys point out the resulting validity problems (Ye et al. 2026). Reading ability from behavior rather than from asking has a precedent in evidence-centered design and stealth assessment (Mislevy, Steinberg, and Almond 2003; Shute and Ventura 2013), where task performance is treated as evidence for a latent competence model. We do not measure the personality of an LLM itself; we measure its ability as a reader where the truth is controlled.

Synthetic players are not new, but their role here is. The procedural personas tradition (Holmgård et al. 2014, 2015,

2019) implements play styles as utility weights and sends evolved personas to play levels in place of human testers, for the purpose of content evaluation. Closest to our work is Shergadwala, Teng, and Seif El-Nasr (2021), which implements behavior tendencies as parameters of a cognitive decision model and recovers them by inverse Bayesian inference. We differ on three points. Personas serve content evaluation and, to our knowledge, have not been used as a ground-truth benchmark for trait-recovery systems. Both lines treat their parameters as ground truth by definition, whereas we admit a parameter only after it passes a behavioral consistency test. And their inference holds the generative model itself, a form of same-model inversion, whereas our readers have no access to the generating policy. Generating behavior from known parameters and testing recovery is otherwise a standard structure (Ng and Russell 2000; Baker, Saxe, and Tenenbaum 2009; Talts et al. 2020); what is new is its organization into a protocol where the ground truth must first be admitted and failures can be localized.

Player modeling and content generation. Both have been studied for decades (Yannakakis and Togelius 2018; Shaker, Togelius, and Nelson 2016; Bunian et al. 2017), and dynamic difficulty adjustment has long closed the loop between them (Hunicke 2005; Mortazavi, Moradi, and Vahabie 2024; Lopes, Fachada, and Fonseca 2025). What the loop has never had is a way to check its own first step: recent surveys note that player modeling is nearly absent from LLM research (Gallotta et al. 2024), and the LLM work that does touch games goes straight to generation (Sudhakaran et al. 2023; Todd et al. 2023). Our generation side is deliberately simple, because the contribution is not a new generator but a reading stage whose output can be validated before it drives one.

Problem and Pipeline

The setup is a recovery problem. A synthetic player whose traits we set plays a level; the session is written down in a form that keeps both what the player did and what was available to do; a LLM reader tries to recover the traits from that record alone, and because we set the traits, its estimate can be checked. The rest of this section states this precisely.

We characterize a player by a latent vector of behavioral tendencies. At each decision moment t during play, the player faces a state s_t and a set of currently feasible choices A_t , and takes an action a_t . Behavior is generated by an unknown policy:

$$a_t \sim \pi(\cdot \mid s_t, A_t; \theta), \quad \theta \in [0, 1]^4.$$

The central point is that observed behavior is determined jointly by tendency and opportunity: the same θ produces different behavior under different A_t . If a record keeps only what was done and discards what could have been done, tendency and opportunity become inseparable. A session is therefore recorded as a sequence of opportunity-aware decision moments:

$$\tau = \{(s_t, A_t, a_t)\}_{t=1}^T, \quad \hat{\theta} = f(\tau),$$

where f is the reader under evaluation, instantiated here mainly as a large language model. Recovery quality is mea-




Figure 1: Overview. (A) Synthetic players with parameterized traits produce opportunity-aware records; parameters enter the benchmark only after passing the admission test, and trait recovery is scored against them. (B) The same record schema drives the deployed pipeline: calibration play, LLM reading, and profile-conditioned generation.

sured by per-dimension rank correlation $\rho(\hat{\theta}_k, \theta_k)$ and per-dimension error. Downstream, $\hat{\theta}$ maps through a fixed function to a target difficulty $D^* = g(\hat{\theta})$, and the generator brings the assembled level close to that target.

The pipeline runs calibration play, opportunity-aware recording, trait reading, difficulty mapping, and profile-conditioned generation, with replay feeding back into the record (Figure 1B). Each stage is checked against controlled ground truth, which is where the four research questions come from.

Method

The four parts answer requirements that follow from one another. Traits must be explicit parameters entering separable channels, or no reader could tell them apart (Synthetic Players). A parameter is not a valid trait until it survives a test on behavior it did not define (Ground-Truth Admission). A reader can only recover what the record preserves, and behavior without its choice set is ambiguous (Records and Reading). And an inferred profile is only useful if it changes the level the player receives (Profile-Conditioned Adaptation).

Synthetic Players

We need players whose traits are caused rather than labeled: behavior must follow from a parameter we set, not from an

annotation assigned afterward. Each synthetic player is a real-time decision maker in a bullet-hell shooter running at sixty frames per second, with the decision module decoupled from the rendering engine and controlled by a single random seed, so that the trait parameters and the seed reproduce the complete trajectory frame by frame.

The policy has three components, kept separate so that the traits do not interfere: if two traits shared a channel, no reader could separate them, and the admission test below would reject both. Objective arbitration decides when to leave avoidance and pursue an objective, with pursuit intent resampled every two seconds from the preference parameters; this stochastic expression proved superior to persistent lock-on, which decouples the parameter from behavior. The movement layer scores candidate positions by a danger field and vetoes directions failing a time-to-collision gate. The action layer evaluates discrete abilities at fixed intervals, first judging whether an opportunity exists. Full definitions are in the appendix.

Trait parameterization. Each of the four trait parameters $\theta_k \in [0, 1]$ enters exactly one channel: Dodge the movement layer, Collector and Aggression the arbitration module, Skill the action layer. The preference dimensions enter through an intent gate, pursuing when $u_t < \theta_k$ with $u_t \sim U(0, 1)$ resampled every two seconds, so the expected share of pursuit time equals θ_k ; both also shrink the safety margin required before pursuit, $T_{\text{safe}}(\theta_k) = T_0(1 - 0.5\theta_k)$, and Aggression addition-

ally controls press depth beneath enemies. The Dodge dimension linearly scales four movement-layer coefficients (reaction interval, lookahead, perception radius, motor noise), of which the reaction interval is the most consequential:

$$\Delta t(\theta_{\text{dodge}}) = 420 - 340 \theta_{\text{dodge}} \text{ ms.} \quad (1)$$

The endpoints bracket the human range rather than imitate it: 420 ms is about twice the mean simple visual reaction time of a 1469-adult sample (Woods et al. 2015), and 80 ms is below any human latency, so the sweep spans impaired to superhuman rather than compressing six levels into the human band. Between two decisions the agent executes its stale target blindly, so a slow reactor walks into bullets that appear inside the decision gap. The Skill dimension is the most direct channel, a conditional use probability $P(\text{use} \mid \text{opportunity}) = \theta_{\text{skill}}$, where the thresholds defining an opportunity are fixed and do not vary with the parameter.

Ground-Truth Admission

Writing a parameter into code does not make it ground truth: it may fail to produce the intended behavioral effect, or may be entangled with other parameters. We therefore subject every parameter to an admission test before it enters the benchmark. Let m_k denote the signature behavioral metric of dimension k , chosen to be *held out*: it plays no role in defining the parameter itself. Each parameter is swept over $L = 6$ levels ($\theta_k \in \{0, 0.2, \dots, 1.0\}$) with five random seeds per level ($n = 30$ runs per dimension), and admitted only if three conditions hold together, instantiating construct validation (Cronbach and Meehl 1955) in a setting where the construct is known by construction. **C1, monotonicity** (convergent validity): the parameter correlates monotonically with m_k across levels, at Spearman $\rho \geq 0.6$ with no significant non-monotone segment. **C2, isolation** (discriminant validity (Campbell and Fiske 1959)): varying the parameter induces no systematic drift in the signature metrics of the other dimensions. **C3, separability**: between-level differences exceed seed-induced variance, quantified by an effect-size ratio and adjacent-level discriminability. Exact statistics and thresholds for C2 and C3 are in the appendix. Configurations that fail do not enter the benchmark and are reported as rejected or non-identifiable; we froze all metrics, directions, and thresholds before observing any model response.

Opportunity-Aware Records and Trait Reading

We serialize each session into a plain-text record whose schema is the same for every level, so that content regenerated per player still yields the same kind of readable object. The design principle is to keep opportunities and choices on file together: recording only what the player *did* is insufficient for judging what the player *wanted*, since the count of opportunities is the denominator against which willingness is judged. What the record does *not* contain is the outcome an unselected action would have produced; the representation is choice-set-aware, not counterfactual.

A record contains decision-moment lines with declined opportunities alongside taken ones, combat events, win-dowed summaries of positioning and opportunity density,

and a closing tally with the fixed game constants needed as denominators. An excerpt appears in the appendix, including a collection opportunity the agent declines, on file precisely so that *did not* can be told apart from *could not*. Serialization follows two disciplines: no trait or metric name ever appears in a record, enforced by a blacklist check; and the wording follows the empirical distribution of behavior rather than the visual layout of the screen.

Because personalized generation means that no two players see the same content, a reader fitted to one level’s statistics would need refitting for every level it produces, and no labels exist for content that does not exist until it is generated. The reader is therefore prompted rather than trained. It operates in a single few-shot setting (Brown et al. 2020): the system prompt contains the game rules, behavioral definitions of the four dimensions, and scoring guidance grounded in quantities countable from the record; the context adds a table of roughly one hundred rows mapping behavioral metrics to verified parameters, and the two most similar sessions retrieved by behavioral features, each with its verified profile (Liu et al. 2022). The example library and the 90-session test set are disjoint at the run level, with the query session and any run sharing its trait configuration excluded from retrieval.¹ Each session is read three times and the scores averaged, with the session as the unit of inference.

Profile-Conditioned Adaptation

An inferred profile is worth nothing unless it changes what the player receives, and unless that change can be checked by something the generator does not control. Dodge and Aggression drive difficulty: a response surface fitted on measured data maps them to a target difficulty, which is converted into a pair of enemy-fire settings and applied uniformly at level assembly, deterministically, so the same profile locks to the same difficulty across repeated generations. Collector, Skill, and a flanking facet of Aggression are expressed through a content style layer that adjusts reward placement, ability-opportunity cadence, and enemy approach angles; the templates are in the appendix.

To separate *matching* a player from merely *easing* the game, we define a difficulty coefficient

$$D_L = \sigma\left(\sum_i w_i \ln \frac{x_i}{x_i^{\text{ref}}}\right) \in [0, 1], \quad (2)$$

the sigmoid of a weighted sum of log ratios of level content features x_i against a standard calibration level, anchored so that the standard level sits at $D_L = 0.5$. The weights w_i are fixed a priori from measured feature elasticities, fire pressure being the dominant lever, rather than learned, so that D_L is an interpretable content-pressure coefficient rather than a fitted success predictor. On this coefficient a novice receives a level with $D^* \approx 0.13$ and an expert receives $D^* \approx 0.57$: matching goes in both directions, while easing goes in only one. D_L controls generation only; downstream evaluation uses independent behavioral outcomes, so the system is not measured with its own ruler.

¹The reader does not update model parameters but does consume labeled data in context, so we avoid describing the setting as zero-training.

Those outcomes are scored on a flow-matching scale (Csikszentmihalyi 1990; Chen 2007) where lower is better: a run whose hits per minute falls inside a band frozen in advance from mid-level agents scores zero, deviation outside the band scores up to one, and a run ending in death scores above one, ranked by how early it ended. The jump at one is deliberate, since hits per minute cannot distinguish a death from a short clean run. The band is a fixed experiential target, so adaptation is judged by whether it brings agents of any level into it, not by whether each level meets a level-specific criterion. The piecewise definition is in the appendix.

Experiments

Setup

Data. Ground-truth admission uses 30 runs per dimension (six levels, five seeds each). Recovery uses 210 labeled training sessions and a frozen test set of 90. The identifiability analysis uses 300 randomly parameterized sessions, the opportunity-context ablation 300 paired sessions, and the diagnostic probe 30 sessions per condition.

LLMs and Baselines. We evaluate two LLM readers, GPT-5.6-sol (OpenAI 2026) and Qwen3.5-122B-A10B-think (Qwen Team 2026), and compare them with rule-based, supervised, retrieval-based, and statistical baselines. Rule uses handcrafted behavioral signatures. FeatReg applies RBF kernel ridge regression (Hoerl and Kennard 1970; Schölkopf and Smola 2002) over behavioral features (a linear-ridge variant scores slightly lower); EmbReg applies kernel ridge over transcript embeddings from the OpenAI text-embedding-3-small model,² an approach reported to work well for text-based trait inference elsewhere (Maharjan et al. 2025), which makes its performance on our records (RQ2) informative rather than a weak-baseline artifact. BPM 1-NN matches each session to its nearest labeled behavioral prototype (Cover and Hart 1967). The soft and hard GMM baselines (Dempster, Laird, and Rubin 1977) estimate traits using eight behavioral clusters.

Implementation. Prompt construction, retrieval exclusion, and three-read averaging follow the Method section. We report per-dimension Spearman ρ and MAE with macro averages.

RQ1: When Does a Parameter Qualify as Ground Truth

Before asking whether a reader can recover a trait, we have to ask whether the trait is there to be recovered. A parameter that produces no consistent behavioral signature is not ground truth, whatever the code says. All four dimensions pass the admission threshold with monotone responses (Figure 2), but their distinguishable resolutions differ markedly. This is the point of the admission test: it exposes weak constructs rather than hiding them.

We further analyze multidimensional identifiability on 300 sessions with random parameters, defining the behavioral-




Figure 2: Trait-parameter monotonicity validation: each parameter produces the expected monotone trend in its held-out behavioral probe (mean $\pm$ std over 5 seeds per level, 6 levels each).

twin collision ratio

$$r_k = \frac{\mathbb{E}[|\Delta\theta_k| \mid \text{behavioral nearest-neighbor pairs}]}{\mathbb{E}[|\Delta\theta_k| \mid \text{random pairs}]}, \quad (3)$$

where a ratio near 1 means that individuals with the same behavior are unconstrained in that dimension, i.e., observationally equivalent. Skill is easiest to identify ($r = .51$) and Collector hardest ($r = .82$): two agents with nearly identical behavior can hold collector parameters far apart. The feature signal of Aggression is in fact ample—positional evidence shows the strongest effect among all metrics—so its recovery difficulty stems mostly from the reading stage rather than from a lack of signal in the environment.

That a linear regressor recovers these parameters from a few hundred labeled sessions is expected rather than surprising: the admission test selects precisely for parameters whose behavioral effect is consistent, monotone, and separable. A clean feature-to-parameter mapping is what an admitted parameter means; parameters without one were rejected before the benchmark was formed.

RQ2: Can External Readers Recover the Parameters

With ground truth admitted, the question becomes what can recover it from behavior alone, and how much of any shortfall belongs to the reader rather than to the record. On a frozen test set of 90 sessions (a preregistered split), all methods read the same complete records and are scored on the same partition; Table 1 reports Spearman ρ per dimension.

Three observations follow. First, feature regression is best read as an instrument rather than a competitor. It answers a different question: given labels drawn from the same distribution as the test set, how much trait information do these records contain at all? Its macro ρ of .77 answers that they

²<https://platform.openai.com/docs/guides/embeddings>

Method	Dodge	Coll.	Aggr.	Skill	Macro
Rule	.67	<u>.50</u>	.45	.89	.63
FeatReg	.77	.56	.83	.91	.77
EmbReg	.65	-.04	.23	.67	.38
BPM 1-NN	.53	.32	.58	.84	.56
GMM (soft, $K=8$)	.36	.38	.62	.76	.53
GMM (hard)	.32	.38	.56	.75	.50
GPT-5.6-sol	.69	.43	<u>.79</u>	<u>.90</u>	<u>.70</u>
Qwen3.5-122B	.62	.49	.69	.82	.65

Table 1: Recovery on the frozen test partition (Spearman ρ). Best in **bold**, second best underlined. Supervised baselines are trained on 210 labeled sessions drawn from the same distribution as the test set; the LLM rows are our few-shot protocol instantiated with the indicated backbone.

contain a great deal, which is what makes every other reader’s shortfall interpretable; had it reached only .40, the failure would lie in the representation and no reader could do better. The supervised baselines are trained on 210 labeled sessions and evaluated on the held-out 90, both partitions drawn from the same distribution, so this instrument presumes a labeled corpus in-distribution with the test set, a condition that cannot hold for real players whose true traits are unobservable.

Second, the LLM reader with GPT-5.6-sol attains .70 macro, second overall, while consuming only a fixed, offline example library and no gradient updates; the same protocol on Qwen3.5-122B-A10B-think reaches .65, above the rule baseline, so the protocol is not tied to a single backbone. Third, Collector is the hardest dimension for every method, consistent with its collision ratio in Eq. 3: the difficulty is in the environment’s expressiveness, not in any particular reader.

The embedding regression reads the same transcript text but pools it into a single fixed-length vector. The opportunity-choice structure that distinguishes a declined pickup from an absent one is a local, countable relation between specific lines, and mean-pooling averages it away; the collapse is sharpest on Collector ($\rho \approx 0$), whose signal lives almost entirely in that relation, while dimensions whose signal survives as global statistics (Dodge, Skill) retain moderate correlation.

Where the reader loses, and why. The gap to feature regression is uneven across dimensions, and its shape identifies its cause. On Skill the two are level (.91 versus .90); the evidence is a countable frequency that survives serialization intact. Aggression is nearly level as well (.83 versus .79). The gap on Dodge (.77 versus .69) tracks quantization: grazing distance and reaction latency are continuous quantities that the transcript reports in coarse categories, which a regressor reads directly from the numeric feature. The largest gap is on Collector (.56 versus .43), and it is environmental rather than architectural: on the diagnostic probe level, where collection opportunities actually arise, the same reader attains $\rho = .92$ (Table 3), matching the handcrafted signature and exceeding linear ridge regression. What the standard level denies the reader is opportunity, not capability.

RQ3: Does Failure Lie in the Environment

A recovery score that falls short says nothing about where to look. Two explanations are always available: the record may not carry the evidence, or the environment may never have produced it. This section tests both. Table 2 reports the opportunity-context ablation; these runs use a different configuration, so absolute values differ from Table 1.

	Dodge	Coll.	Aggr.	Skill
ρ full	.61	.39	.50	.74
ρ ablated	.61	.35	.55	.59

Table 2: Opportunity-context ablation (Spearman ρ , 300 paired runs).

Opportunity information carries weight. Removing what *could have been done* from the records degrades precisely the two dimensions that depend on opportunities for expression—Skill falls from .74 to .59 and Collector from .39 to .35—while Dodge is unchanged. Aggression, by contrast, improves slightly (.50 to .55), suggesting that the opportunity-related wording was not helping the reader on this dimension. The opportunity-aware representation is not a redundant design.

An environment defect, not a reader defect. Table 1 raises a question it cannot answer on its own: why is Collector the weakest dimension for every reader, handcrafted metrics included? The cause turns out to lie upstream of reading. In standard levels, reward objects drop from kills, so an agent that fights little almost never faces the choice of whether to collect; a strong collector and an indifferent one behave nearly alike, because neither is ever offered the choice. After constructing a diagnostic level that actively supplies reward objects, Collector recoverability rises from .43 to .92 on LLM readings, and Table 3 shows the improvement is not specific to any reader: every family reads Collector at $\rho \geq .84$ there. To read out a trait, the environment must first give it an opportunity for expression, which is also why content that adapts to the player is not merely a convenience: a fixed level measures only the traits it happens to afford.

Reader	Type	ρ	MAE
Handcrafted signature (pursue rate)	metric	.92	—
Rule (calibrated signature)	sup. (LOO)	.91	.097
Linear ridge (17 features)	sup. (LOO)	.91	.121
RBF kernel ridge	sup. (LOO)	.84	.165
BPM prototype (1-NN)	retrieval	.85	.153
kNN (5-NN)	retrieval	.88	.165
LLM (GPT-5.6-sol)	LLM	.92	.144
LLM (Qwen3.5-122B-A10B-think)	LLM	.89	.137

Table 3: Collector recovery on the diagnostic probe level (Spearman ρ , $n = 30$ per condition).

RQ4: From Recovery to Adaptation

Adaptation could improve outcomes for the wrong reason: any level made easier will kill fewer players, whether or not the profile driving it is correct. The five arms exist to separate those two explanations. The downstream experiment has five arms: adaptation from ground truth, adaptation from the LLM-inferred profile, uniform difficulty, a mismatched profile, and plain easing. Each arm is run once per agent level on a single generated level, and scored on the flow-matching scale defined above.

Bot	Truth	Inferred	Mismatch	Fixed	Easing
Novice	1.04 [†]	1.21 [†]	1.17 [†]	1.25 [†]	1.07 [†]
Low-mid	0	0	1.15 [†]	1.19 [†]	1.16 [†]
Mid	0	0.33	1.06 [†]	1.02 [†]	0
High	0	0	1.13 [†]	0	0
Expert	0	0.19	1.19 [†]	0.21	0.28
Survived	4/5	4/5	0/5	2/5	3/5

Table 4: Flow-matching score per agent level and arm (lower is better; **0** = inside the flow band). [†] marks a run that ended in death, which scores above 1 by construction. One run per cell.

With one run per cell, the decimals do not support fine comparison, but the survival row is unambiguous. The mismatched profile fails at every level, which is the control that matters: adaptation depends on the profile being right, not merely on the game being easier. Ground truth and the inferred profile both survive four of five runs, and they differ only in degree, the inferred profile landing outside the flow band at the mid and expert levels where ground truth lands inside it. That residual gap is the downstream cost of reading error, and it is smaller than the gap between either of them and the profile-free arms. Uniform difficulty and plain easing survive two and three runs, and both fail at the two lowest levels. The novice level defeats every arm, including ground truth: the lowest difficulty the generator can assemble is still above what that agent survives, which is a limit of the generation side rather than of the reading.

Human Study

Everything so far is synthetic by construction, which is what makes it checkable and also what limits it. The question the title asks is whether reading behavior beats asking, when the player is a person. Twelve participants first played a fixed baseline level and then completed a self-report questionnaire on the experience (appendix). The questionnaire response and the behavioral record from the same session were used in parallel to assemble two personalized levels, which each participant played without being told which was which. Our LLM player model then read the behavioral records from all sessions. Figure 3 plots the resulting ability readings.

Readings under the questionnaire-generated condition scatter in both directions relative to baseline, several participants falling below it, whereas the behavior-generated condition is higher for every participant.




Figure 3: Pilot study ($n = 12$). Per-participant LLM ability readings under three level-generation conditions; grey lines connect the same participant.

Limitations

This paper builds the apparatus that behavior-based player modeling has lacked and then uses it. A controlled synthetic population makes four traits true by construction, and an admission test decides which of them counts as ground truth at all. An opportunity-aware record keeps declined choices beside taken ones, and ablating the choice sets degrades exactly the traits whose expression depends on opportunity. On this foundation, an LLM player model reads the record with a fixed example library and no gradient updates, recovering traits close to a supervised ceiling that presumes labels real players cannot supply. Closing the loop, the inferred profile conditions a generator, and in a pilot with twelve participants, levels built from behavior scored higher than levels built from a questionnaire.

Three boundaries mark where these results stop. The readers do not see identical inputs: feature regression consumes continuous per-session statistics, whereas the reader sees several of them quantized into bands, so part of the remaining gap belongs to the transcript rather than to the reading, and a representation that preserved those quantities is the obvious next thing to test. The profile is read once from a calibration session and then held fixed, so nothing here speaks to players whose traits shift as they learn the game; the replay loop the pipeline provides is the natural place to track that, and we do not evaluate it. And the synthetic environment trades the noise of human behavior for control: it is the trade that makes every claim above checkable, and it is the reason twelve participants are directional rather than conclusive. Yet every gap is traced to a cause, and an evaluation that can localize failure comes closer to the meaning of validation than a demonstration that can only display success.

References

Baker, C. L.; Saxe, R.; and Tenenbaum, J. B. 2009. Action Understanding as Inverse Planning. *Cognition*, 113: 329–

- Brown, T. B.; et al. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Bunian, S.; Canossa, A.; Colvin, R.; and Seif El-Nasr, M. 2017. Modeling Individual Differences in Game Behavior Using HMM. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment (AIIDE)*, volume 13, 158–164.
- Campbell, D. T.; and Fiske, D. W. 1959. Convergent and Discriminant Validation by the Multitrait-Multimethod Matrix. *Psychological Bulletin*, 56(2): 81–105.
- Chen, J. 2007. Flow in Games (and Everything Else). *Communications of the ACM*, 50(4): 31–34.
- Cover, T. M.; and Hart, P. E. 1967. Nearest Neighbor Pattern Classification. *IEEE Transactions on Information Theory*, 13(1): 21–27.
- Cronbach, L. J.; and Meehl, P. E. 1955. Construct Validity in Psychological Tests. *Psychological Bulletin*, 52(4): 281–302.
- Crowne, D. P.; and Marlowe, D. 1960. A New Scale of Social Desirability Independent of Psychopathology. *Journal of Consulting Psychology*, 24(4): 349–354.
- Csikszentmihalyi, M. 1990. *Flow: The Psychology of Optimal Experience*. Harper and Row.
- Dempster, A. P.; Laird, N. M.; and Rubin, D. B. 1977. Maximum Likelihood from Incomplete Data via the EM Algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)*, 39(1): 1–22.
- Gallotta, R.; Todd, G.; Zammit, M.; Earle, S.; Liapis, A.; Togelius, J.; and Yannakakis, G. N. 2024. Large Language Models and Games: A Survey and Roadmap. *IEEE Transactions on Games*. ArXiv:2402.18659.
- Goel, D.; Mahmoudi-Nejad, A.; and Guzdial, M. 2024. Label-Free Subjective Player Experience Modelling via Let's Play Videos. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, volume 20, 46–53.
- Hoerl, A. E.; and Kennard, R. W. 1970. Ridge Regression: Biased Estimation for Nonorthogonal Problems. *Technometrics*, 12(1): 55–67.
- Holmgård, C.; Green, M. C.; Liapis, A.; and Togelius, J. 2019. Automated Playtesting with Procedural Personas through MCTS with Evolved Heuristics. *IEEE Transactions on Games*, 11(4): 352–362.
- Holmgård, C.; Liapis, A.; Togelius, J.; and Yannakakis, G. N. 2014. Evolving Personas for Player Decision Modeling. In *Proceedings of the 2014 IEEE Conference on Computational Intelligence and Games (CIG)*, 1–8.
- Holmgård, C.; Liapis, A.; Togelius, J.; and Yannakakis, G. N. 2015. Monte-Carlo Tree Search for Persona Based Player Modeling. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment (AIIDE)*, volume 11, 8–14.
- Hunicke, R. 2005. The Case for Dynamic Difficulty Adjustment in Games. In *Proceedings of the 2005 ACM SIGCHI International Conference on Advances in Computer Entertainment Technology (ACE)*, 429–433.
- John, O. P.; and Srivastava, S. 1999. The Big-Five Trait Taxonomy: History, Measurement, and Theoretical Perspectives. In *Handbook of Personality: Theory and Research*, 102–138. Guilford Press, 2nd edition.
- Liu, J.; Shen, D.; Zhang, Y.; Dolan, B.; Carin, L.; and Chen, W. 2022. What Makes Good In-Context Examples for GPT-3? In *Proceedings of Deep Learning Inside Out (DeeLIO), ACL Workshop*.
- Lopes, P.; Fachada, N.; and Fonseca, M. 2025. Closing the Loop: A Systematic Review of Experience-Driven Game Adaptation. ArXiv:2505.01351.
- Maharjan, J.; Jin, R.; Zhu, J.; and Kenne, D. 2025. Psychometric Evaluation of Large Language Model Embeddings for Personality Trait Prediction. *Journal of Medical Internet Research*, 27: e75347.
- Mislevy, R. J.; Steinberg, L. S.; and Almond, R. G. 2003. On the Structure of Educational Assessments. *Measurement: Interdisciplinary Research and Perspectives*, 1(1): 3–62.
- Mortazavi, F.; Moradi, H.; and Vahabie, A.-H. 2024. Dynamic Difficulty Adjustment Approaches in Video Games: A Systematic Literature Review. *Multimedia Tools and Applications*, 83.
- Ng, A. Y.; and Russell, S. 2000. Algorithms for Inverse Reinforcement Learning. In *Proceedings of the Seventeenth International Conference on Machine Learning (ICML)*, 663–670.
- Nisbett, R. E.; and Wilson, T. D. 1977. Telling More Than We Can Know: Verbal Reports on Mental Processes. *Psychological Review*, 84(3): 231–259.
- OpenAI. 2026. GPT-5.6: Frontier Intelligence That Scales with Your Ambition.
- Peters, H.; Cerf, M.; and Matz, S. C. 2024. Large Language Models Can Infer Personality from Free-Form User Interactions. arXiv:2405.13052.
- Podsakoff, P. M.; MacKenzie, S. B.; Lee, J.-Y.; and Podsakoff, N. P. 2003. Common Method Biases in Behavioral Research: A Critical Review of the Literature and Recommended Remedies. *Journal of Applied Psychology*, 88(5): 879–903.
- Qwen Team. 2026. Qwen3.5-122B-A10B.
- Schölkopf, B.; and Smola, A. J. 2002. *Learning with Kernels: Support Vector Machines, Regularization, Optimization, and Beyond*. Cambridge, MA: MIT Press.
- Shaker, N.; Togelius, J.; and Nelson, M. J. 2016. *Procedural Content Generation in Games*. Springer.
- Shergadwala, M. N.; Teng, Z.; and Seif El-Nasr, M. 2021. Can We Infer Player Behavior Tendencies from a Player's Decision-Making Data? Integrating Theory of Mind to Player Modeling. In *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment (AIIDE)*, volume 17, 195–202.
- Shute, V.; and Ventura, M. 2013. *Stealth Assessment: Measuring and Supporting Learning in Video Games*. MIT Press.

- Sudhakaran, S.; González-Duque, M.; Freiburger, M.; Glanois, C.; Najarro, E.; and Risi, S. 2023. MarioGPT: Open-Ended Text2Level Generation through Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*. ArXiv:2302.05981.
- Talts, S.; Betancourt, M.; Simpson, D.; Vehtari, A.; and Gelman, A. 2020. Validating Bayesian Inference Algorithms with Simulation-Based Calibration. arXiv:1804.06788.
- Tett, R. P.; and Burnett, D. D. 2003. A Personality Trait-Based Interactionist Model of Job Performance. *Journal of Applied Psychology*, 88(3): 500–517.
- Todd, G.; Earle, S.; Nasir, M. U.; Green, M. C.; and Togelius, J. 2023. Level Generation Through Large Language Models. In *Proceedings of the 18th International Conference on the Foundations of Digital Games (FDG)*.
- Woods, D. L.; Wyma, J. M.; Yund, E. W.; Herron, T. J.; and Reed, B. 2015. Factors Influencing the Latency of Simple Reaction Time. *Frontiers in Human Neuroscience*, 9: 131.
- Yannakakis, G. N.; Spronck, P.; Loiacono, D.; and André, E. 2013. Player Modeling. In *Artificial and Computational Intelligence in Games*, volume 6 of *Dagstuhl Follow-Ups*, 45–59.
- Yannakakis, G. N.; and Togelius, J. 2011. Experience-Driven Procedural Content Generation. *IEEE Transactions on Affective Computing*, 2(3): 147–161.
- Yannakakis, G. N.; and Togelius, J. 2018. *Artificial Intelligence and Games*. Springer.
- Ye, H.; Jin, J.; Xie, Y.; Zhang, X.; and Song, G. 2026. Large Language Model Psychometrics: A Systematic Review of Evaluation, Validation, and Enhancement. arXiv:2505.08245.
- Zhu, J.; Jin, R.; and Coifman, K. G. 2025. Can LLMs Infer Personality from Real World Conversations? arXiv:2507.14355.
- Zhu, J.; Maharjan, J.; Li, X.; Coifman, K. G.; and Jin, R. 2025. Evaluating LLM Alignment on Personality Inference from Real-World Interview Data. arXiv:2509.13244.

A Synthetic-Player Policy

The synthetic player (Bot) is a real-time decision maker running at sixty frames per second, decoupled from rendering and driven by a single random seed: given the four trait parameters $\theta \in [0, 1]^4$ and the seed, the trajectory is reproducible frame by frame. Ground truth is therefore the generative parameter, not a post-hoc annotation. The policy has three layers: objective arbitration, a movement layer, and an action layer.

A.1 Trait Parameterization (Constants)

Each θ_k enters exactly one channel: Dodge scales the movement layer, Collector and Aggression enter the arbitration intent gate, and Skill sets the action layer’s per-opportunity use probability $P(\text{use} \mid \text{opportunity}) = \theta_{\text{skill}}$.

Both preference dimensions enter through the intent gate,

$$\text{pursue}_k \iff u_t < \theta_k, \quad u_t \sim U(0, 1) \text{ resampled every 2 s,} \quad (\text{A1})$$

so the expected share of pursuit time equals θ_k . The interval is re-rolled every 2 s rather than held fixed for a whole pursuit, which keeps a high-aggression bot from locking onto one tanky enemy and washing out the signal. Both also shrink the safety margin required before pursuit,

$$T_{\text{safe}}(\theta_k) = T_0 (1 - 0.5 \theta_k), \quad (\text{A2})$$

and the reaction interval is the dominant Dodge channel, $\Delta t(\theta_{\text{dodge}}) = 420 - 340 \theta_{\text{dodge}}$ ms.

A.2 Movement Layer: Danger Field and Time-to-Collision

At each decision step the agent samples candidate destinations and scores each with a *danger field*, moving toward the reachable point of lowest score subject to a time-to-collision (TTC) veto. Both quantities are computed analytically, so a fast hazard cannot slip between time samples.

Danger field. For candidate point p and hazard b with position–velocity $(\mathbf{d}, \mathbf{v})$, $\mathbf{d} = (b_x - p_x, b_y - p_y)$, the time of closest approach over the foresight horizon H_d is

$$t^* = \text{clip}\left(-\frac{\mathbf{d} \cdot \mathbf{v}}{\|\mathbf{v}\|^2}, 0, H_d\right), \quad (\text{A3})$$

the closest-approach clearance is $m_b = \|\mathbf{d} + t^* \mathbf{v}\| - r_b$, and a hazard contributes only if $m_b < R_d$ (danger radius $R_d = 110$ px):

$$c_b = \frac{R_d - m_b}{R_d} \quad (\text{unclamped for } m_b < 0). \quad (\text{A4})$$

Leaving c_b unclamped gives a continuous outward gradient out of a hazard body. The field is *max-dominated*, not summed:

$$\text{Danger}(p) = c_{\text{max}} + 0.15 \left(\sum_b c_b - c_{\text{max}} \right), \quad c_{\text{max}} = \max_b c_b, \quad (\text{A5})$$

so the acute-threat gradient survives at any bullet density; a pure sum saturates under dense fire and the agent freezes.

Algorithm A1 Synthetic-player decision loop (one frame)

Require: traits θ , state s_t , seed-driven rng

- 1: **if** $t - t_{\text{roll}} \geq 2000$ ms **then**
 - 2: $u_{\text{coll}}, u_{\text{aggr}} \leftarrow \text{rng}()$; $t_{\text{roll}} \leftarrow t$
 - 3: **end if**
 - 4: $\text{obj} \leftarrow$ earliest-deadline objective with $u < \theta$ (Eq. A1), gated by T_{safe} (Eq. A2)
 - 5: sample candidate points; score each by $\text{Danger}(\cdot)$ (Eq. A5)
 - 6: $\text{dir} \leftarrow$ arg min danger among directions with $\text{TTC} \geq 0.18$ s (Eq. A6); else max-TTC
 - 7: **if** ability opportunity present **and** $\text{rng}() < \theta_{\text{skill}}$ **then**
 - 8: fire missile / activate berserk
 - 9: **end if**
 - 10: move toward obj if chosen, else toward dir
-

Time to collision. For the hero at relative position $\mathbf{r}_p$ and velocity $\mathbf{r}_v$ against combined radius r , TTC is the smaller nonnegative root of $\|\mathbf{r}_p + t \mathbf{r}_v\| = r$:

$$\text{TTC} = \frac{-b - \sqrt{b^2 - 4ac}}{2a}, \quad a = \|\mathbf{r}_v\|^2, \quad b = 2 \mathbf{r}_p \mathbf{r}_v, \quad c = \|\mathbf{r}_p\|^2 - r^2, \quad (\text{A6})$$

returning $+\infty$ when the hazard moves apart or does not intersect. A candidate direction whose soonest bullet-or-boundary collision falls below $\text{TTC}_{\text{min}} = 0.18$ s is vetoed; if all directions are unsafe, the agent takes the one with the largest TTC. Hero and enemy collision radii are $0.175 w_{\text{hero}}$ and $0.4 w_{\text{enemy}} + 0.175 w_{\text{hero}}$ respectively ($w_{\text{hero}} \approx 68$ px).

A.3 Decision Loop (Pseudocode)

A.4 Alternative Policy Family

A second *utility family* replaces the heuristics above with utility scoring over candidate actions; the same parameters produce similar tendencies. A full cross-family recovery evaluation is left to future work.

B Ground-Truth Admission

B.1 Signature Metrics

Each dimension is validated against a held-out *signature metric* that plays no role in defining the parameter: Dodge $\rightarrow$ wave-phase hits per exposed minute (expect $\downarrow$); Collector $\rightarrow$ XP pursue rate = picked/spawned (expect $\uparrow$); Aggression $\rightarrow$ mean flight- y position (expect $\downarrow$; lower y means the ship pressed forward toward the enemies—the strongest single feature effect in the benchmark, $\rho \approx -.91$); Skill $\rightarrow$ ability attempts per exposed (non-berserk) minute (expect $\uparrow$).

B.2 Criteria and Thresholds

Each parameter is swept over $L = 6$ levels $\theta \in \{0, .2, .4, .6, .8, 1.0\}$ with 5 seeds per level ($n = 30$ per dimension). All metrics, directions, and thresholds were frozen before any model response was observed. A parameter is admitted iff:

- C1 Monotonicity (convergent validity).** Spearman ρ between level and its signature metric satisfies $\rho \geq 0.6$ with the expected sign and no significant non-monotone segment.
- C2 Isolation (no systematic drift).** The full cross-dimension matrix of Spearman correlations between the swept level and *every* axis’s signature metric is computed; the diagonal must dominate, and any off-diagonal $|\rho| > 0.5$ is flagged as cross-talk.
- C3 Separability.** Adjacent levels, plus the (0.4, 1.0) and (0.6, 1.0) pairs, are compared by an exact Mann–Whitney permutation test (5 vs. 5, all 252 splits) with a directed $AUC = U/(n_1 n_2)$; a pair with $AUC < 0.7$ is flagged indistinguishable and greedily merged into effective levels. A one-way ANOVA F and η^2 across level groups are reported as signal-to-noise summaries but are *not* used as an admission gate.

Configurations that fail are reported as rejected or non-identifiable.

C Behavioral Records

C.1 Format

Each session serializes to plain text containing: decision-moment lines (state + the action taken, with declined opportunities on file alongside taken ones), combat-event lines, ten-second windowed summaries (spawn/kill/escape counts, orbs spawned/picked/expired, grazes, positional band), and a closing whole-session tally. The summaries supply the *denominators* against which willingness is judged (e.g., *orbs spawned* vs. *picked*).

C.2 Leakage Control

No trait name, internal metric key, provenance field, or precomputed statistic ever appears in a record. This is enforced by a denylist check, `assertNoLeakage`, run after every serialization; a hit raises a hard error. The denylist covers (i) trait/axis names (`dodge_skill`, `collector`, `aggression`, `skill_usage`, ...); (ii) metric keys from the measurement pipeline (`hits_rate`, `pursue_rate`, `kill_conversion`, ...); (iii) provenance/labels (`oracle`, `ground truth`, `generated_spec`, `calibration`); (iv) statistics machinery (`z-score`, `sigmoid`, `percentile`, `baseline`); and (v) the bot’s internal decision flags (`wantKill`, `wantOrb`, `safetyGateOpen`). Because a bot session and a human session are serialized by the same procedure into the same leakage-free schema, the reader operates on the same kind of object regardless of origin; the parameter θ generates behavior upstream and never enters what is read.

C.3 Excerpt

The lines below are verbatim from a real *human* session (alpha level, 123.7 s, outcome WIN):

```
[t=23.6s] DECISION engage: hp=175
bombs=0
  nearestBullet=116px danger=0
enemies=3
```

```
berserkActive=no -> EVADED
[t=25.2s] BERSERK activated (hp=175,
  bullets onscreen=16, nearest=60px,
enemies=6)
[t=29.7s] DECISION xp-orb: hp=175
bombs=0
  nearestBullet=143px danger=0
enemies=5
  berserkActive=yes -> AVOIDED
[t=30.0s] SUMMARY 20-30s: enemies
spawned=11 killed=5
  escaped=4 | xp orbs spawned=5
  picked=2 | grazes=7
[t=35.7s] DECISION xp-orb: hp=175
bombs=0
  nearestBullet=173px danger=0
enemies=8
  berserkActive=no -> AVOIDED
```

The two `xp-orb` $\rightarrow$ AVOIDED lines are collection opportunities that were present but declined; on file alongside the `spawned=5` `picked=2` denominator, “did not” can be told apart from “could not”.

D The Reader

D.1 Prompt

The reader is a single fixed function from a transcript to a four-dimensional estimate. It is a prompt, not a trained model: the same prompt is used for every session and every backbone, and it is frozen before any result is read. The full text is reproduced verbatim below. It has three parts: a description of the game and the reading task, a set of scoring rules that fix how observations map to numbers, and the required response format. The transcript is appended after it as the user message.

You are an expert analyst of video-game player behavior. You will receive a behavioral transcript of one run of a vertical-scrolling shoot-em-up ("`shmup`") played by an unknown player, and you must profile that player.

THE GAME: The player pilots a ship over a vertically scrolling battlefield, auto-firing upward at enemies that enter from the top. Enemies shoot bullets; getting hit costs HP and the run ends in defeat at 0 HP. Destroyed enemies drop XP orbs that disappear if not collected; flying close to an orb collects it. Gift drops give weapon upgrades (blue) or missile refills (red). The player has two limited abilities: MISSILES (a screen-clearing burst, limited stock) and BERSERK (a short invulnerable power window on a cooldown, triggered at will). A "graze" is a bullet passing very close without hitting. The run ends by defeating the boss (WIN), dying (LOSS), or hitting a session time cap (TIMEOUT).

THE TRANSCRIPT contains: a header (level, duration, outcome, final

HP, missiles left); DECISION lines capturing the full game state at notable choice points and the action the player took; combat beats (hits taken, missile launches, berserk activations, pickups, boss phases); and per-10-second SUMMARY lines (enemies spawned/killed/escaped, xp orbs spawned/picked/expired, grazes) plus POSITION lines (vertical band the ship flew in -- FRONT is up near the enemies, BACK is the safe bottom -- and how close enemy bullets got).

YOUR TASK: estimate four independent traits of this player. Each is a number in [0,1] where 0 means the extreme low end of the player population and 1 the extreme high end. Judge each trait purely from observable behavior in the transcript:

- "dodge_skill" -- how skillfully the player avoids getting hit while under fire. High: stays unhit even when many bullets are on screen and passing close; keeps control in dense patterns; survives long stretches without losing HP. Low: repeatedly takes hits whenever bullets are around, bleeds HP quickly under pressure. Judge relative to the danger actually faced.

- "collector" -- the player's drive to gather rewards. High: chases down XP orbs and gift drops before they expire, detours toward pickups even from a comfortable position, ends up collecting most of what dropped. Low: ignores orbs and gifts, lets most rewards drift away uncollected.

- "aggression" -- the player's drive to seek out and destroy enemies. High: pushes toward enemies, holds position to finish kills, keeps fighting rather than backing off, lets few enemies escape. Low: hangs back passively, avoids confrontation, lets many enemies leave alive.

- "skill_usage" -- the player's willingness to actually spend the limited abilities (missiles and berserk) when threats or opportunities arise, rather than hoarding them. High: uses berserk and missiles readily and repeatedly. Low: finishes the run with abilities unspent even in dangerous moments.

These traits are independent. Base everything ONLY on the transcript; do not assume a trait value because of the outcome alone.

SCORING RULES -- your numbers must follow these frequency definitions, not a general impression:

- "skill_usage": the FRACTION of clear

opportunities the player actually spent an ability on -- never the raw count of uses. Two denominators are visible: (1) DECISION bomb lines record each missile offer and whether it was USED or HELD -- the USED share is the most direct read; (2) berserk recharges ~10s start-to-start, so ~6 activations fit in a minute of sustained combat: ~5-6/min = 0.9-1.0, ~3/min = 0.4-0.5, 1-2/min while enemies stay onscreen = 0.0-0.3. If the screen was mostly empty, judge from the missile DECISION share alone.

- "aggression": judge from NORMAL-flight positioning (POSITION lines) and engagement: parked deep in BACK band letting most enemies leave = 0.0-0.2; MID band, engaging targets that come close = ~0.5; pressing toward FRONT and hunting most enemies down = 0.8-1.0. Kills made while invincible say little about aggression.

- "dodge_skill": judge from the RATE of HIT lines RELATIVE TO bullet exposure, never from the final HP or whether the run was a LOSS. Taking few hits across many windows of close, dense fire = 0.8-1.0; getting hit in most exposed windows = 0.0-0.2; roughly one hit per couple of minutes of real pressure = ~0.5. Do NOT lower this just because HP ran low or the run ended in defeat. Judge "collector" from its definition above as usual.

Respond with a single JSON object, nothing else:

```
{"dodge_skill": n, "collector": n, "aggression": n, "skill_usage": n, "rationale": "2-4 sentences citing concrete transcript evidence"}
```

Two clauses added for the human study. The human study adds two clauses to the skill-usage rule. First, a DECISION bomb line with bombs=0 is not a declined offer—there was no missile to spend—so it is left out of the USED-share denominator; running the stock down to zero by spending counts *for* high skill usage, not against it. Second, the berserk rate is computed directly (activations ÷ minutes) and the scale applied from that number, rather than judged by impression.

D.2 Few-Shot Retrieval

Two kinds of labeled evidence are placed in context: a table of $N=100$ rows mapping behavioral metrics to verified parameters, and the $k=2$ most similar sessions retrieved by z-scored Euclidean distance over behavioral features, each paired with its verified profile. Two leakage guards apply: the query run itself (matched by run identifier) is never retrieved, and any exact feature duplicate (distance ≈ 0) is dropped. All labeled evidence originates from a fixed calibration level, so labeling is paid once.

D.3 Decoding and Repeats

The reader uses the provider defaults (no explicit temperature or top- p) with a JSON-object response format; each session is read 3 times independently and the per-dimension scores averaged, with the session as the unit of statistical analysis. The reported backbones are gpt-5.6-sol and Qwen3.5-122B-A10B-think.

E Baselines

We compare the reader against five baselines, all evaluated on the same frozen 90-session test set. The supervision available to the two sides is deliberately *not* equal. Every baseline is fitted on the 210 labeled training sessions, drawn from the same distribution as the test set; the reader performs no gradient updates and consumes only a fixed, offline example library placed in context. The supervised baselines are therefore best read as an instrument rather than as competitors: they measure how much trait information the records contain when in-distribution labels are available, a condition that cannot hold for real players whose true traits are unobservable. Four of the five read the same 17 z -scored behavioral features (Table A1); one (EmbReg) reads the transcript text itself, exactly as the reader does.

Rule — handcrafted signatures. The simplest baseline, and the “obvious feature, read directly” reference. Each trait is predicted from one hand-picked statistic—the same kind of held-out signature used to admit that trait in Section B (e.g. hits per exposed minute for Dodge, XP pursue rate for Collector, mean flight- y for Aggression, ability attempts per minute for Skill)—rescaled monotonically to $[0, 1]$. It has no learned parameters.

FeatReg — feature kernel ridge. The strongest baseline overall. It fits one kernel ridge regressor per trait, with an RBF kernel over the 17 features; the ridge penalty λ is chosen per trait by inner 4-fold cross-validation over $\lambda \in \{1, 10, 100, 1000\}$. A plain linear-ridge version is only slightly weaker (macro ρ .75 vs. .77), so most of the accuracy comes from the features rather than the kernel. FeatReg is the ceiling referred to above: given same-distribution labels, it measures how much trait signal the records actually contain.

EmbReg — embedding kernel ridge. The same kernel-ridge setup, but over the transcript *text* instead of the engineered features. Each transcript is embedded with OpenAI text-embedding-3-small (1536-d) and mean-pooled to one vector, then regressed to the four traits with the same λ selection. Because it reads exactly what the reader reads, it isolates whether an off-the-shelf text embedding already carries the trait signal the reader extracts.

BPM 1-NN — behavioral-prototype matching. A non-parametric memory baseline: each test session is matched to its single nearest training session in the z -scored 17-feature space, and takes that neighbor’s four labels as its prediction. It asks whether trait recovery is just “retrieve the most similar labeled player.”

FeatReg feature set (17)

hits_rate, graze_rate, mean_bullet_dist, gun_kill_conversion, kill_conversion, kills_per_min, escape_share, xp_pursue_rate, xp_expired_share, gift_rate, missile_rate, berserk_rate, ability_rate, mean_y, std_y, front_share, invincible_share

Table A1: Behavioral features used by the supervised feature baselines.

Feature	w_i	Feature	w_i
trackerShare	.25	effectiveHp	.10
bulletsPerSec	.20	dangerExposure	.10
spawnRatePerMin	.15	bulletSpeedMul	.00
peakConcurrent	.15	reliefPerMin	−.05

Table A2: D_L prior weights ($\sum |w| = 1$). bulletSpeedMul is measured non-monotone and excluded; reliefPerMin is a relief term (negative).

GMM — soft and hard clustering. Fits eight Gaussian clusters with diagonal covariance to the training features by EM (seed 7, 200 iterations, variance floor 10^{-2}); each cluster stores the mean trait vector of its member sessions. The *soft* variant predicts a responsibility-weighted average of those cluster trait-means; the *hard* variant uses only the single most likely cluster. It tests whether a handful of behavioral “player types” is enough to place a session.

F Difficulty Model

F.1 Content Difficulty Coefficient D_L

$D_L = \sigma(\sum_i w_i \ln(x_i/x_i^{\text{ref}}))$ measures how much pressure a level’s content puts on the player. Each content feature x_i is compared to its value on a standard calibration level as a log ratio $\ln(x_i/x_i^{\text{ref}})$; the weighted sum of those ratios is passed through a sigmoid, anchored so the standard level sits at $D_L = 0.5$. The weights w_i are fixed in advance from measured feature elasticities, not learned, so D_L is an interpretable content-pressure coefficient rather than a fitted success predictor (Table A2).

F.2 Target Surface $g(\cdot)$

The target difficulty $D^* = g(\hat{\theta})$ is set from the estimated Dodge and Aggression. The mapping runs through a response surface hits/min = $f(\text{dodge, fireMul})$ fitted on 167 runs (a dual-difficulty battery); the fire multiplier is clamped to $[0.18, 3.0]$.

F.3 Flow-Matching Outcome Scale

Downstream outcomes are scored on a flow-matching scale (lower is better). The score is read from a behavioral quantity the generator does not control—hits per minute, h . Let $[L, H]$ be a flow band frozen in advance from mid-level agents, C

the session-length cap, and T the run duration:

$$\text{score} = \begin{cases} 1 + \frac{C - T}{C} & \text{death } (T < C, \text{ no win}) \\ \min(1, \frac{h-H}{H}) & \text{survived, } h > H \text{ (overload)} \\ \min(1, \frac{L-h}{L}) & \text{survived, } h < L \text{ (boredom)} \\ 0 & \text{survived, } L \leq h \leq H \text{ (flow)}. \end{cases} \quad (\text{A7})$$

The jump at 1 is deliberate: hits per minute cannot tell a fast death from a short clean run, so a death is placed above any survived outcome and ranked by how early it ended. The band used for the downstream flow-matching score is $[L, H] = [2.8, 7.3]$ hits/min, frozen in advance from mid-level reference agents (mean $5.07 \pm 1.5\sigma$); it is distinct from the generation-side targeter band and is never used to control generation.

G Generation Pipeline

G.1 Director and Translation

A recovered profile becomes a playable level in four stages. The first two are done by a language model (which we call the *director*); the last two are deterministic code.

(1) *Set the target*. The profile is turned into a director input. Alongside the profile, this input states the target difficulty as an average scene intensity—a single number $[\text{Intensity } X/10]$ derived from D^* . Stating the difficulty this way makes the director generate *toward* a difficulty, instead of free-styling from the profile alone.

(2) *Draft the level*. The director (`gpt-5.6-sol`) writes a structured draft: a world, an enemy roster, a boss, and 5–9 *scenes*. Each scene carries an intensity that never drops from one scene to the next, so the level ramps up.

(3) *Translate to waves*. A deterministic step turns each scene’s enemy groups into concrete game waves. It reads the scene’s free-text intent to pick a formation and a movement—“flanking” becomes a pincer, “V wedge” a v-formation—and spreads each group’s `count` out in space. Without this step every wave would collapse into a single column falling straight down.

(4) *Stage trait content*. A second deterministic step (the *rule layer*) adds content tuned to the profile on top of the draft. It adds three things. First, rewards that drop regardless of kills: missile refills, plus—for players read as high collectors—pickups placed in riskier spots. These give the Collector trait a way to show itself that killing enemies cannot; without them a level might never offer a pickup to decline in the first place (the opportunity-scarcity problem of Section C). Second, evenly spaced ability opportunities on the ~ 10 s berserk cooldown, keyed to the Skill estimate, so how often the player spends an ability is countable. Third, a set of numeric difficulty multipliers.

The full director prompt is in the Code and Data Supplement. The director never sees θ or the raw features—only the derived profile and the intensity number—so the ground-truth label cannot leak through the generator.

G.2 Content Style-Layer Templates

The rule layer gives each of the three opportunity-mediated traits its own content template, controlling reward placement, ability-opportunity cadence, and enemy approach angles. Each template is built only from the same wave settings a normal level uses—enemy HP multiplier ($\text{hpMul} \leq 2$), fire-interval multiplier ($\text{fireIntervalMul} \geq 0.5$), enemy counts, and spawn cadence—never from special measurement-only controls. A session played over a template therefore serializes to the same leakage-free transcript as any other level.

- **Collector (reward placement)**. A standing center bullet field (a single strong enemy at $\text{hpMul} = 2.0$, the schema cap, respawning every 2.6 s across $x \in \{140, 200, 260, 320\}$ with $\text{fireIntervalMul} = \text{lerp}(1.4, 0.6, c)$ so field density rises with the Collector estimate c), plus orb offers every 2.5 s whose x is drawn from a safe-edge→center ladder $[80, 400, 120, 360, 160, 320, 200, 280, 240]$. The estimate slides the offer-mix center ($\text{lerp}(2, 8, c)$ over the ladder): low collectors are offered mostly safe-edge pickups, high collectors mostly center (danger-priced) pickups, with a few of each kept as anchors.
- **Aggression (enemy approach angles)**. Sparse single tanky spawns (strong, $\text{hpMul} = 1.6$, forward fire) on a 5 s cadence, entering from an alternating flank cycle $x \in [90, 390, 180, 300, 240]$; each is an independent pursue-or-ignore offer with a straight-down escape route, so pushing forward to engage is a choice the record captures rather than a forced encounter.
- **Skill (ability-opportunity cadence)**. A 12 s countable schedule that alternates a soft cluster ($8 \times \text{dijil}$, berserk/missile bait) with a pressure spike ($3 \times \text{strong}$ in a v-formation, $\text{fireIntervalMul} = 0.6$, panic bait), evenly spaced so spend-per-opportunity is directly readable from the transcript.

Dodge is not templated: it is read from bullet-proximity behavior on whatever field the other templates already produce, so it needs no opportunity scaffold of its own.

G.3 Closed-Loop Difficulty Refinement

The mapping from a profile to its enemy-fire settings is deterministic: the same profile yields the same settings and the same target on every generation. What the refinement loop below adjusts is the assembled *content*, not that mapping, and it is applied only to the synthetic downstream levels; the human study uses the draft directly.

Drafting a level does not pin how hard it actually plays: the director’s sense of “how many enemies make intensity 7” is noisy, so two levels drafted for the same target can play far apart. To pin difficulty for the synthetic downstream experiments, a refinement loop measures each level and edits it. It plays the level with one *fixed* bot—the same bot for every level and target, so any difference in the reading comes from the content and not from a changing player—and reads back the difficulty the level actually produced, written D_L^{real} (the difficulty coefficient D_L of Section F, measured on that play).

Algorithm A2 Closed-loop difficulty refinement

Require: level spec s , target τ , tolerance $\epsilon=0.06$

```
1: density  $\leftarrow 1$ ; escalation  $\leftarrow 0$ 
2: for  $i = 0$  to MAXITERS do
3:    $D \leftarrow$  probe-bot play of  $s$ ; measure  $D_L^{\text{real}}$ 
4:   if  $|D - \tau| \leq \epsilon$  then
5:     return  $s$ 
6:   end if
7:   adjust fire-density knob toward  $\tau$  (proportional, on the
   logit scale)
8:   if knob saturated and  $D \ll \tau$  then
9:     escalate enemy kinds: promote fodder waves (tail-
       first) to elite kinds (tracker fire + hover dwell +
       hp)
10:  else if  $D \gg \tau$  then
11:    de-escalate one level
12:  end if
13: end for
```

It then edits the content until D_L^{real} meets the target (Algorithm A2). The bot’s bias is calibrated once on the standard level ($D_L^{\text{real}}=0.573$, target = $D^*+0.073$).

The lever that actually moves difficulty is upgrading enemy *type*, not adding more enemies. Turning up fire rate and wave count both stall at $D_L^{\text{real}} \approx 0.27$: weak, fast-leaving enemies (*fodder*) and undirected fire rarely come near the ship. Upgrading a wave to a stronger *elite* type changes three things at once—the enemy aims its fire at the ship (*tracker fire*), lingers on screen instead of flying past (*hover dwell*), and takes more hits to kill. The single heaviest term in D_L , the share of aimed-fire enemies (`trackerShare`), is zero on an all-fodder level, which is why *type* is the effective knob. To keep the level’s style, upgrades are applied from the last scenes backward (preserving the ramp) and reuse whatever elite type the level already uses. Across target levels this cuts the mean absolute target error from 0.204 (draft only) to 0.011, usually within 2–3 plays.

H Dataset

The Code and Data Supplement includes the synthetic benchmark: two frozen 300-session populations, each session sampled at an independent random parameter vector under a fixed seed. Every session ships with its ground-truth parameter label (`oracleTraits`), the leakage-free transcript the reader consumes, and a session summary (level, seed, outcome, duration). A `schema.md` documents every field. Because a human’s latent traits are unobservable, no existing public dataset offers verifiable trait labels; this is the gap the synthetic benchmark fills, and the reason recovery accuracy can be measured at all.

I Human Study

I.1 Procedure

Twelve participants each play a fixed baseline (calibration) level and complete a short self-report questionnaire. Each then plays two further personalized levels, in a counterbalanced order they are not told. Both personalized levels are

produced by the same LLM generator (the director) from a four-axis player profile; the two conditions differ only in where that profile comes from. In the **behavior-driven** condition, the profile is the one the LLM infers from the participant’s own play. In the **questionnaire-driven** condition, the profile is read straight off the participant’s questionnaire self-report (Part A) and fed to the same generator in place of the inferred one. Because the generator is held fixed, the comparison isolates the source of the profile—read from behavior versus asked directly. After the sessions, the LLM player model reads the behavioral records, and the ability reading on each personalized level is compared against the baseline reading.

I.2 Questionnaire

The instrument has two parts. Part A is a self-report of play style: its four axis scores form the profile that drives the questionnaire-driven level, and are also compared against the LLM reading. Part B collects background. Responses are on a 7-point Likert scale (1 = strongly disagree, ..., 7 = strongly agree) unless an item says otherwise. Part A is filled immediately after the baseline level and *before* the participant sees any model output, to prevent anchoring; Part B is filled last.

Part A — self-reported play style. Each axis pairs a forward (+) and a reverse (−) item; the axis score is $\text{mean}(\text{forward}, 8 - \text{reverse})$, normalized to $[0, 1]$ by $(\text{score} - 1)/6$. These four scores are the profile passed to the generator in the questionnaire-driven condition.

- A1 (*Dodge*, +) Even when bullets are dense, I can steadily avoid most of them.
- A2 (*Dodge*, −) When I get hit, it is mostly because I could not react in time.
- A3 (*Collector*, +) Whenever a pickup or supply appears, I try to go get it.
- A4 (*Collector*, −) If grabbing a pickup is risky, I would rather skip it.
- A5 (*Aggression*, +) I tend to push to the front of the screen to attack rather than stay in a safe zone.
- A6 (*Aggression*, −) I prioritize survival first, and only then think about kills.
- A7 (*Skill*, +) As soon as an ability (missile / berserk) is ready, I tend to spend it rather than hoard it.
- A8 (*Skill*, −) I often finish a run with abilities still unused.

Part B — background.

- B1 Experience with bullet-hell / flight-shooter games: 1 = never played, ..., 7 = very experienced.
- B2 Average weekly time spent playing games: none / <2h / 2–7h / 7–15h / >15h.

I.3 Analysis

The primary readout plotted in the study figure is the change in the macro ability reading $\hat{\theta}$ (the mean over the four axes) from the baseline level to each personalized level; the comparison of interest is whether the behavior-driven level lifts

$\hat{\theta}$ more than the questionnaire-driven level. As a secondary check, the per-axis Part A self-report is compared with the LLM reading, with the self-report treated as a noisy, potentially circular criterion rather than as accuracy. At $n=12$ these are reported as directional, not significance-tested.